

XMPP for Cloud Services

Cloud Camp – Milan
10 september 2009

Alessandro Malgaroli (Bluendo)

alex@jabber.blundo.com

Fabio Forno (Bluendo)


ff@jabber.blundo.com

Salvatore Loreto (Ericsson Research)

salvatore.loreto@ieee.org

What is Cloud Computing?

- Resources are provided as a service over the Internet
 - Dynamic Scalability
 - High Availability
- Exploding in popularity – a critical trend of software architectures
- Started with simple services; now growing complex
 - Now predominately based on REST architectures, web applications, web services



Cloud Architecture Problems

- REST is exclusively based on the Request/Response pattern
 - No push and asynchronous interactions
 - Polling doesn't scale and isn't real-time;
 - Need of two-way data exchange

(Comet/Bosh/WebSocket could be used to make some improvements)


- Web services (SOAP) are overly complicated when handling:
 - Presence (availability) and discovery of remote services
 - Federation with third party services, access behind NAT and firewalls
 - Many-to-Many distribution pattern, asynchronous or multi step calls
 - Binary data, data streaming

Thesis: *web services are great for simple cloud services or direct p2p interactions; XMPP is better for complex cloud services*

XMPP & Cloud Computing

XMPP with its extensions is a powerful protocol for cloud services that demonstrate several advantages over HTTP-based Web Services:

- Realtime capabilities
 - Eg. Heartbeats, alarms, asynchronous webservices
- Efficient distribution of data: publish/subscribe, direct push
 - Eg. Configuration distribution, push RSS/Atom, data collection, log processing, results delivery to clients
- Federated services & discovery
 - Easy interdomain RPCs
 - Access control based on JIDS and JID domains
 - Advanced discovery capabilities



XMPP Architecture

- XMPP provides a technology for asynchronous, end-to-end exchange of structured data.
- XMPP architectural style builds an overlay network having:
 - Global addresses (JIDs)
 - Network availability (presence)
 - Concurrent information transactions
 - Distributed, federated network
 - Structured Data with XML payloads
- The architecture similar to that of the email network, but it introduces several modifications to facilitate **near-real-time communication**

"Availability for Concurrent Transactions"(ACT)

Peter Saint-Andre


<https://stpeter.im/index.php/2009/09/01/active-architectures>

XMPP Architecture

- Global Addresses:
 - As with email, XMPP uses globally-unique addresses.
 - All XMPP entities are addressable on the network:
- clients
- servers
- additional services accessible by clients and servers.
- Presence:
 - the ability for an entity to advertise its network availability or "presence" to other entities via dedicated communication primitives (the <presence/> stanza).

XMPP Architecture

- Distributed Network, comprising servers, clients and components (server extensions)
- End-to-end communication in XMPP is logically peer-to-peer but physically client-to-server-to-server-to-client
- Federation across different domains



XMPP Architecture

- Data is sent through persistent **XML streams**
- The basic unit of meaning in XMPP is an “**XML stanza**”:
 - a fragment of XML that is sent over a stream
 - the root element of a stanza includes routing attributes (such as "from" and "to" addresses),
 - there are three basic stanzas with different behavior
 - the child elements of the stanza contain an extensible payload for delivery to the intended recipient

XMPP Stanzas

| | Pattern | Examples |
|-----------------|------------------------------------|--|
| Presence | Broadcast to subscribed JIDs | <pre><presence from='juliet@capulet.org'> <c xmlns='http://jabber.org/protocol/caps' .../> </presence></pre> |
| Message | One way Asynchronous | <pre><message from='romeo@montague.org' to='juliet@capulet.org'> <event xmlns='http://jabber.org/protocol/pubsub#event'> ... </event> </message></pre> |
| IQ | Request Response | <pre><iq from='romeo@montague.org' to='balcony.capulet.org' type='get' id='123'> <query xmlns='jabber:iq:disco#items' /></iq> <iq from='balcony.capulet.org' to='romeo@montague.org' type='result' id='123'> <query xmlns='jabber:iq:disco#items'><items>...<items> </query></iq></pre> |

An example: Publish/Subscribe

- Observer pattern:
 - Data is not directly sent to consumers
 - Producers **publish** data to an intermediate service, into topics or nodes
 - Consumers **subscribe** to intermediate nodes
 - Pubsub usually implemented as a dedicated component in servers, any JID can be subscriber or publisher
- Applications
 - Log monitoring, data collection and processing, atom, *wave style applications*
- Why not (only) JMS?
 - **Global naming / Federation / Presence / Sophisticated access model**

PubSub & REST

- Each node is a **resource** where it is possible to do actions and it is identified by URIs:
 - xmpp:pubsub.acme.org;/atom/sport
 - xmpp:user@example.org;/pep/location
- Each node contains items, which are **stateful representation** of the resource
- Each node support **actions**:
 - PUBLISH: publish (PUT) an item or update (POST) an item
 - DELETE: retract an item
 - GET: retrieve a published item
 - LIST: retrieve the available items
 - **SUBSCRIBE: be notified of any change in node items**
- **REST + Realtime notifications + many-to-many**
 - Delivery of notifications is bound to presence (no messages lost)

PubSub Applications

■ Implementations

- Builtin in most servers: tigase, ejabberd, openfire
- Separate components: idavoll (<http://idavoll.ik.nu/>)
- Clients: psi, lampiro, iphone3g, buddycloud, ...
- Libs: wokkel (python), smack (limited support, java)
- Application servers: DEM (<http://freedom.sf.net/>)

